

FAST TAG FRAUD DETECTION

USING MACHINE LEARNING FOR REAL-TIME FRAUD DETECTION

SONAL SONARGHARE



INTRODUCTION

OVERVIEW:

- FASTag is an electronic toll collection system in India.
- Fraud detection is crucial to prevent revenue loss.
- This project applies machine learning to detect fraudulent FASTag transactions.



PROBLEM STATEMENT

PROBLEM:

- Increasing fraudulent FASTag transactions.
- Manual detection methods are inefficient.
- Significant revenue losses due to undetected fraud.

OBJECTIVE:

- Develop a machine learning solution for real-time fraud detection in FASTag transactions.



DATASET DESCRIPTION

DATASET DESCRIPTION:

1. **Transaction_ID**: Unique identifier for each transaction.
2. **Timestamp**: Date and time of the transaction.
3. **Vehicle_Type**: Type of vehicle involved in the transaction.
4. **FastagID**: Unique identifier for Fastag.
5. **TollBoothID**: Identifier for the toll booth.
6. **Lane_Type**: Type of lane used for the transaction.
7. **Vehicle_Dimensions**: Dimensions of the vehicle.
8. **Transaction_Amount**: Amount associated with the transaction.
9. **Amount_paid**: Amount paid for the transaction.
10. **Geographical_Location**: Location details of the transaction.
11. **Vehicle_Speed**: Speed of the vehicle during the transaction.
12. **Vehicle_Plate_Number**: License plate number of the vehicle.
13. **Fraud_indicator**: Binary indicator of fraudulent activity (target variable).



EXPLORATORY DATA ANALYSIS (EDA)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_curve, auc
from imblearn.over_sampling import SMOTE
import warnings
warnings.filterwarnings('ignore')
```

← Import Python Libraries

```
file_path = 'FastagFraudDetection.csv'
# Read the CSV file into a pandas DataFrame
df = pd.read_csv(file_path)
```

← Read CSV File

EXPLORATORY DATA ANALYSIS (EDA)

df.head()

	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensions	Transaction_Amount	Amount_paid	Geographical_Location	Vehicle_Speed	Vehicle_Plate_Number	Fraud_indicator
0	1	1/6/2023 11:20	Bus	FTG-001-ABC-121	A-101	Express	Large	350	120	13.059816123454882, 77.77068662374292	65	KA11AB1234	Fraud
1	2	1/7/2023 14:55	Car	FTG-002-XYZ-451	B-102	Regular	Small	120	100	13.059816123454882, 77.77068662374292	78	KA66CD5678	Fraud
2	3	1/8/2023 18:25	Motorcycle	NaN	D-104	Regular	Small	0	0	13.059816123454882, 77.77068662374292	53	KA88EF9012	Not Fraud

df.tail()

	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensions	Transaction_Amount	Amount_paid	Geographical_Location	Vehicle_Speed	Vehicle_Plate_Number	Fraud_indicator
4995	4996	1/1/2023 22:18	Truck	FTG-445-EDC-765	C-103	Regular	Large	330	330	13.21331620748757, 77.55413526894684	81	KA74ST0123	Not Fraud
4996	4997	1/17/2023 13:43	Van	FTG-446-LMK-432	B-102	Express	Medium	125	125	13.21331620748757, 77.55413526894684	64	KA38UV3456	Not Fraud
4997	4998	2/5/2023 5:08	Sedan	FTG-447-PIN-109	A-101	Regular	Medium	115	115	13.21331620748757, 77.55413526894684	93	KA33WX6789	Not Fraud

df.shape
(5000, 13)

Display basic information about the dataset

```
df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Transaction_ID  5000 non-null   int64  
 1   Timestamp        5000 non-null   object  
 2   Vehicle_Type    5000 non-null   object  
 3   FastagID         4451 non-null   object  
 4   TollBoothID     5000 non-null   object  
 5   Lane_Type        5000 non-null   object  
 6   Vehicle_Dimensions  5000 non-null   object  
 7   Transaction_Amount  5000 non-null   int64  
 8   Amount_paid      5000 non-null   int64  
 9   Geographical_Location  5000 non-null   object  
 10  Vehicle_Speed    5000 non-null   int64  
 11  Vehicle_Plate_Number  5000 non-null   object  
 12  Fraud_indicator  5000 non-null   object  
dtypes: int64(4), object(9)
memory usage: 507.9+ KB
```

Shape & Info →

df.describe()

	Transaction_ID	Transaction_Amount	Amount_paid	Vehicle_Speed
count	5000.000000	5000.000000	5000.000000	5000.000000
mean	2500.500000	161.062000	141.261000	67.851200
std	1443.520003	112.44995	106.480996	16.597547
min	1.000000	0.000000	0.000000	10.000000
25%	1250.750000	100.000000	90.000000	54.000000
50%	2500.500000	130.000000	120.000000	67.000000
75%	3750.250000	290.000000	160.000000	82.000000
max	5000.000000	350.000000	350.000000	118.000000

First few Rows

Last few Rows

Summary Statistics

EXPLORATORY DATA ANALYSIS (EDA)

...

Check for Unique Values

```
df.nunique()
```

Transaction_ID	5000
Timestamp	4423
Vehicle_Type	7
FastagID	4451
TollBoothID	6
Lane_Type	2
Vehicle_Dimensions	3
Transaction_Amount	20
Amount_paid	23
Geographical_Location	5
Vehicle_Speed	85
Vehicle_Plate_Number	5000
Fraud_indicator	2
dtype: int64	

Check for Missing Values

```
df.isnull().sum()
```

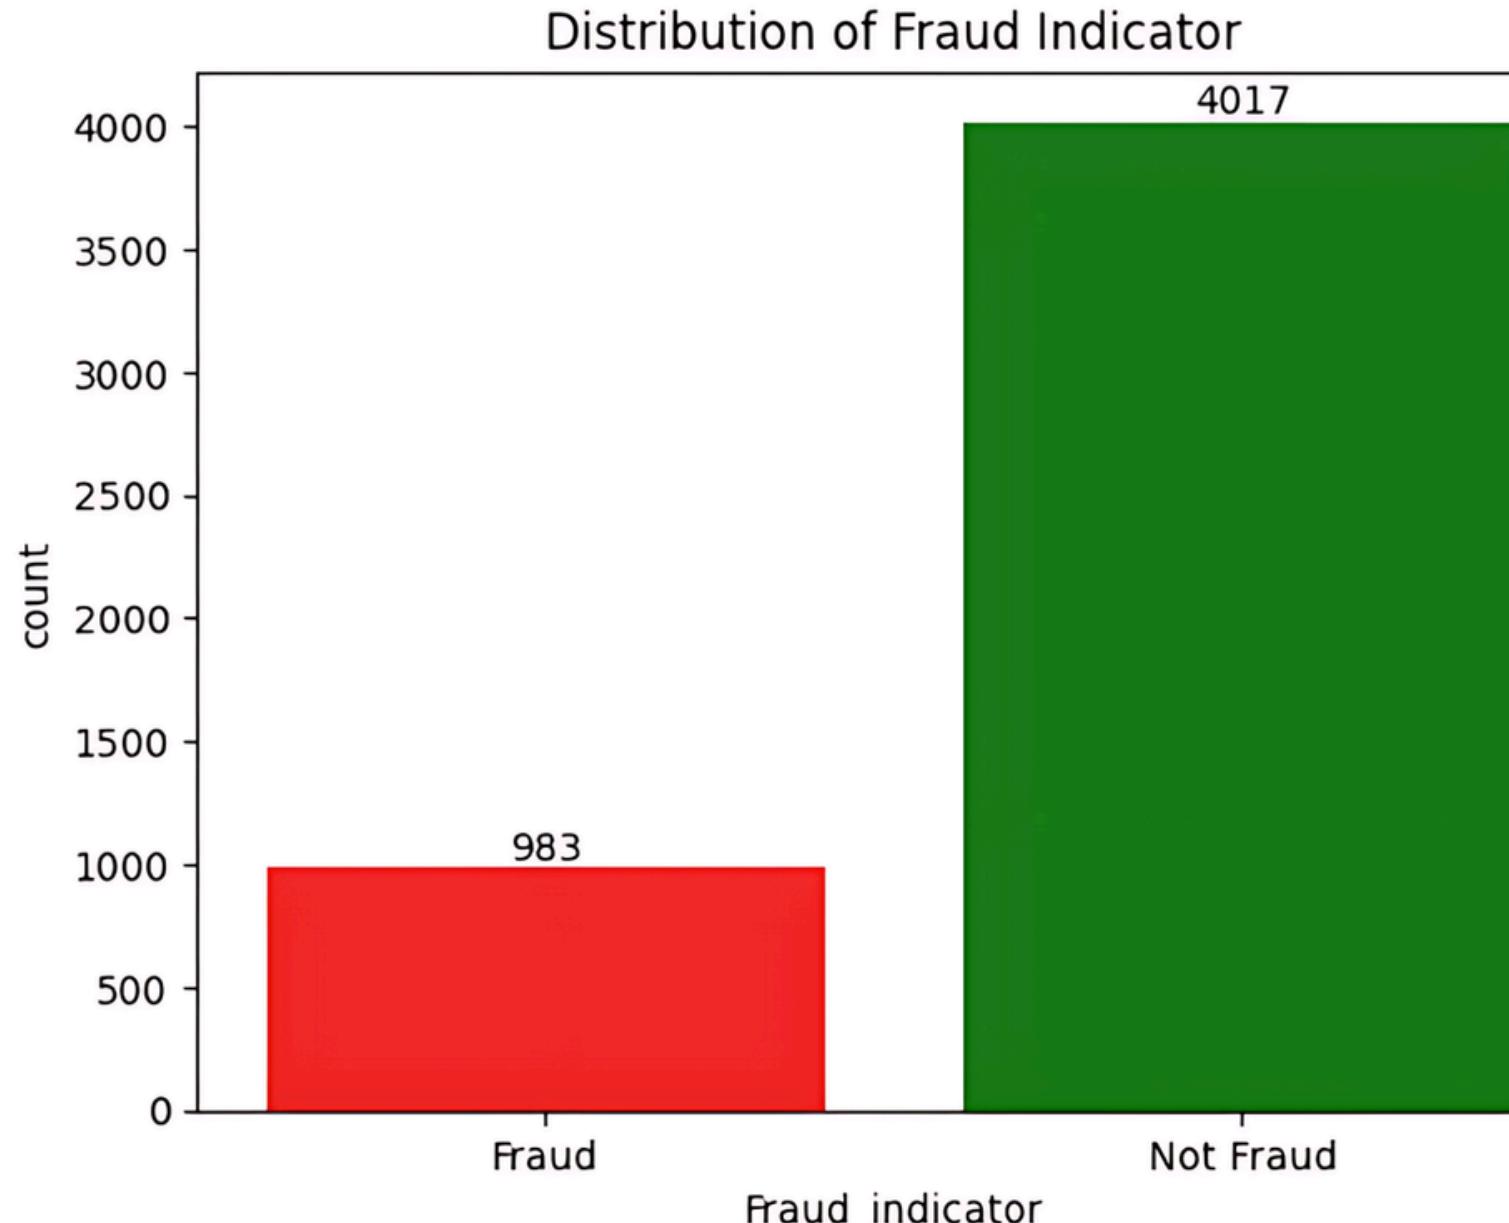
Transaction_ID	0
Timestamp	0
Vehicle_Type	0
FastagID	549
TollBoothID	0
Lane_Type	0
Vehicle_Dimensions	0
Transaction_Amount	0
Amount_paid	0
Geographical_Location	0
Vehicle_Speed	0
Vehicle_Plate_Number	0
Fraud_indicator	0
dtype: int64	

...

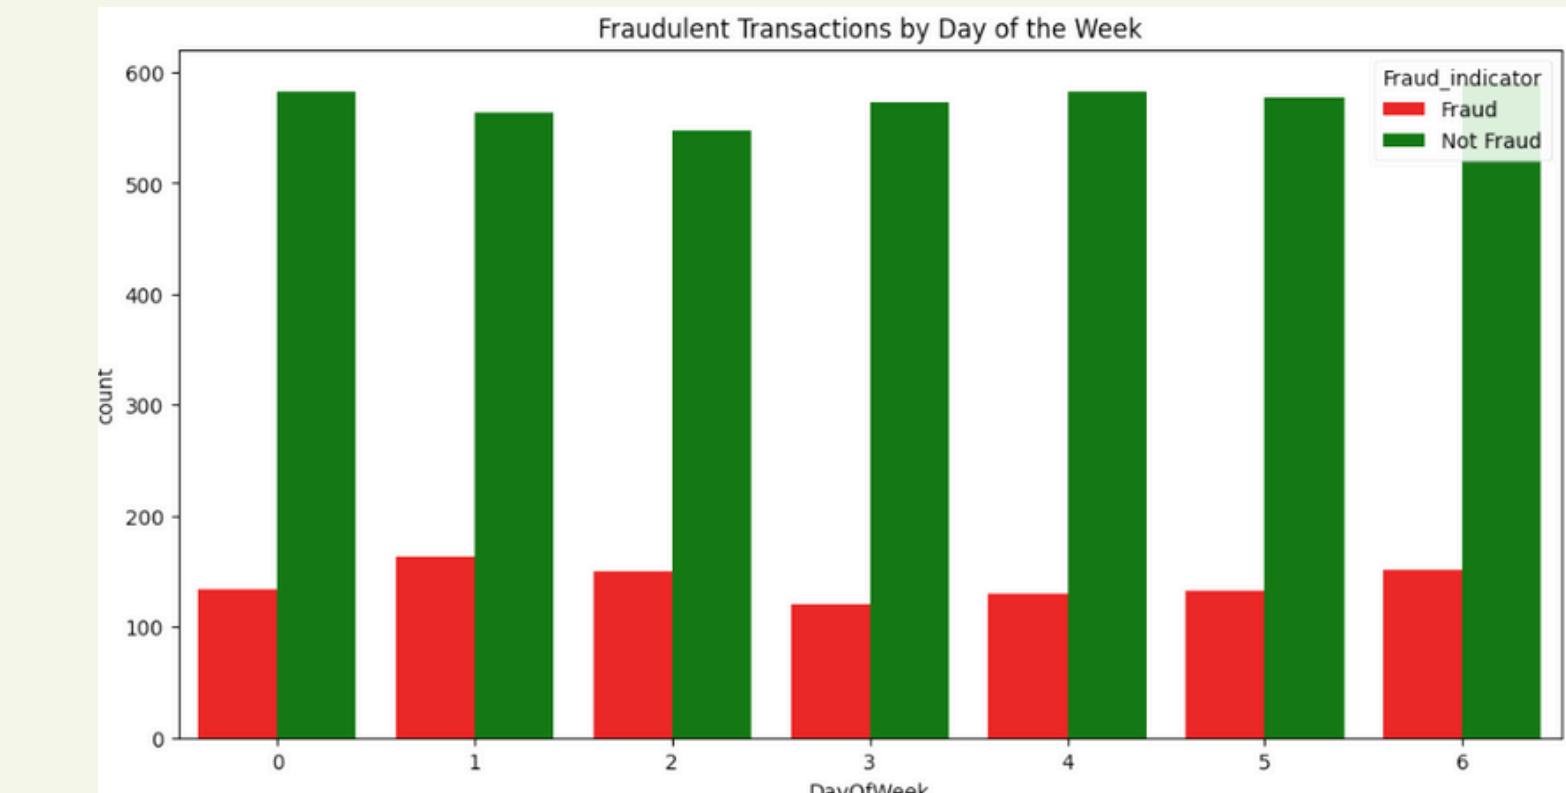
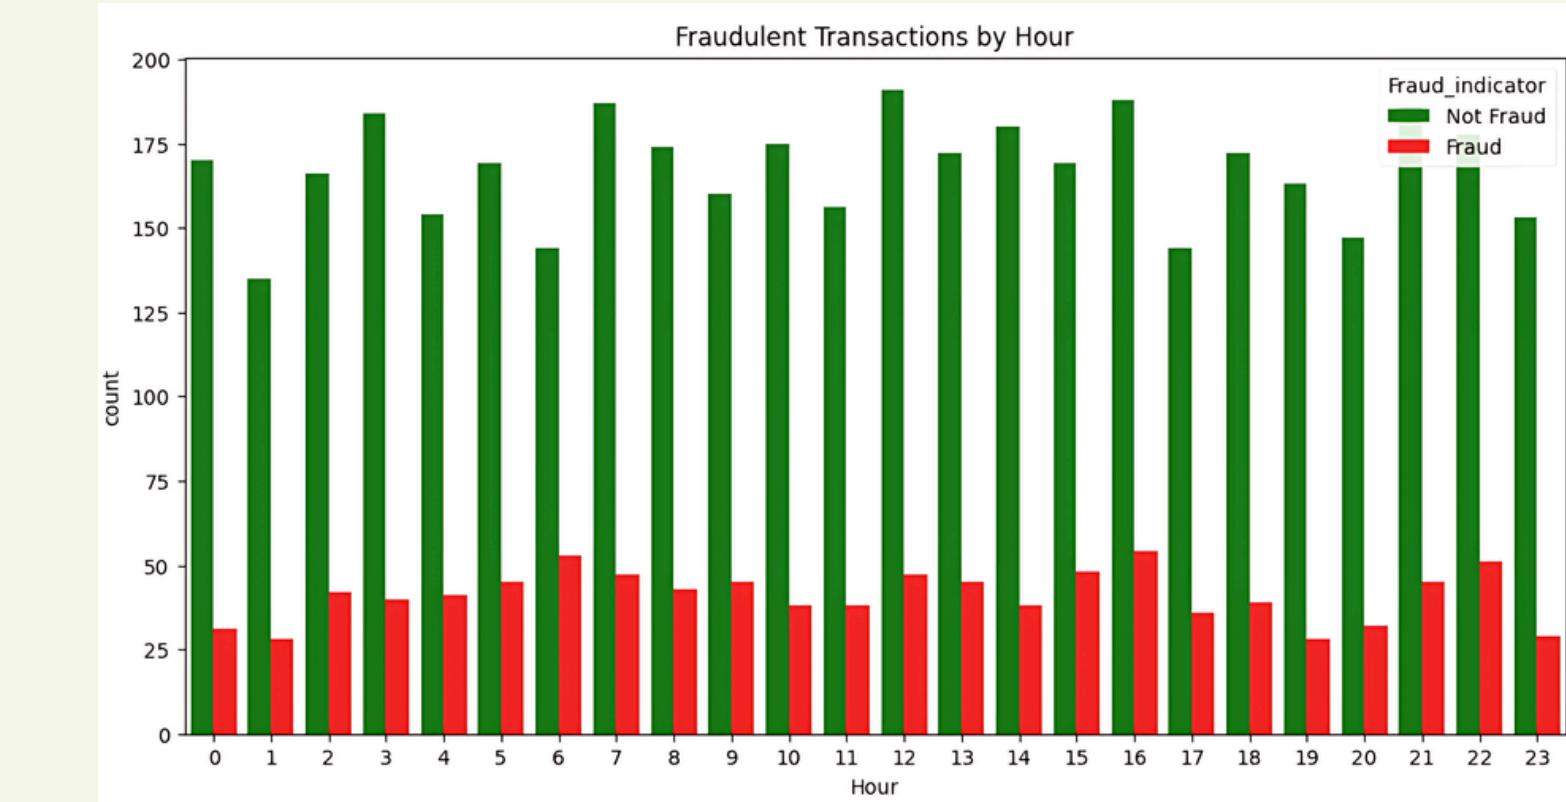
DATA VISUALIZATION

FRAUD DISTRIBUTION

```
# Create the bar plot
colors = {'Fraud': 'red', 'Not Fraud': 'Green'}
ax = sns.countplot(data=df, x='Fraud_indicator', palette=colors)
ax.set_title('Distribution of Fraud Indicator')
for bars in ax.containers:
    ax.bar_label(bars)
```

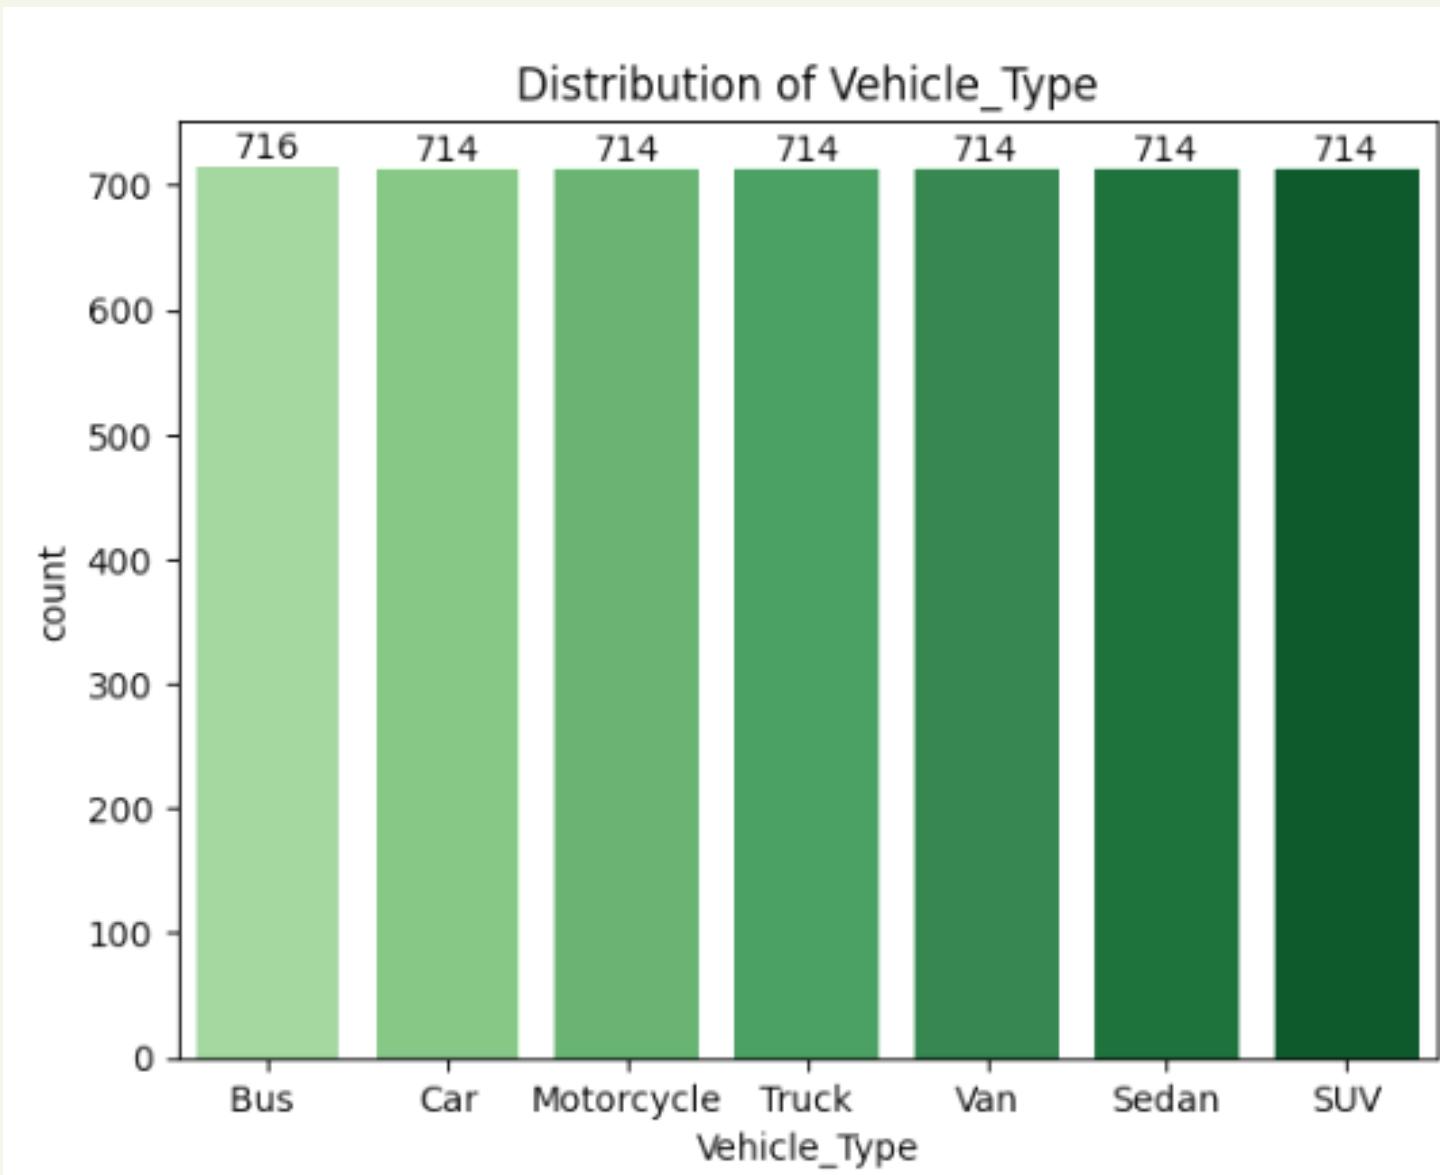


Fraudulent transactions by hour & DayOfWeek

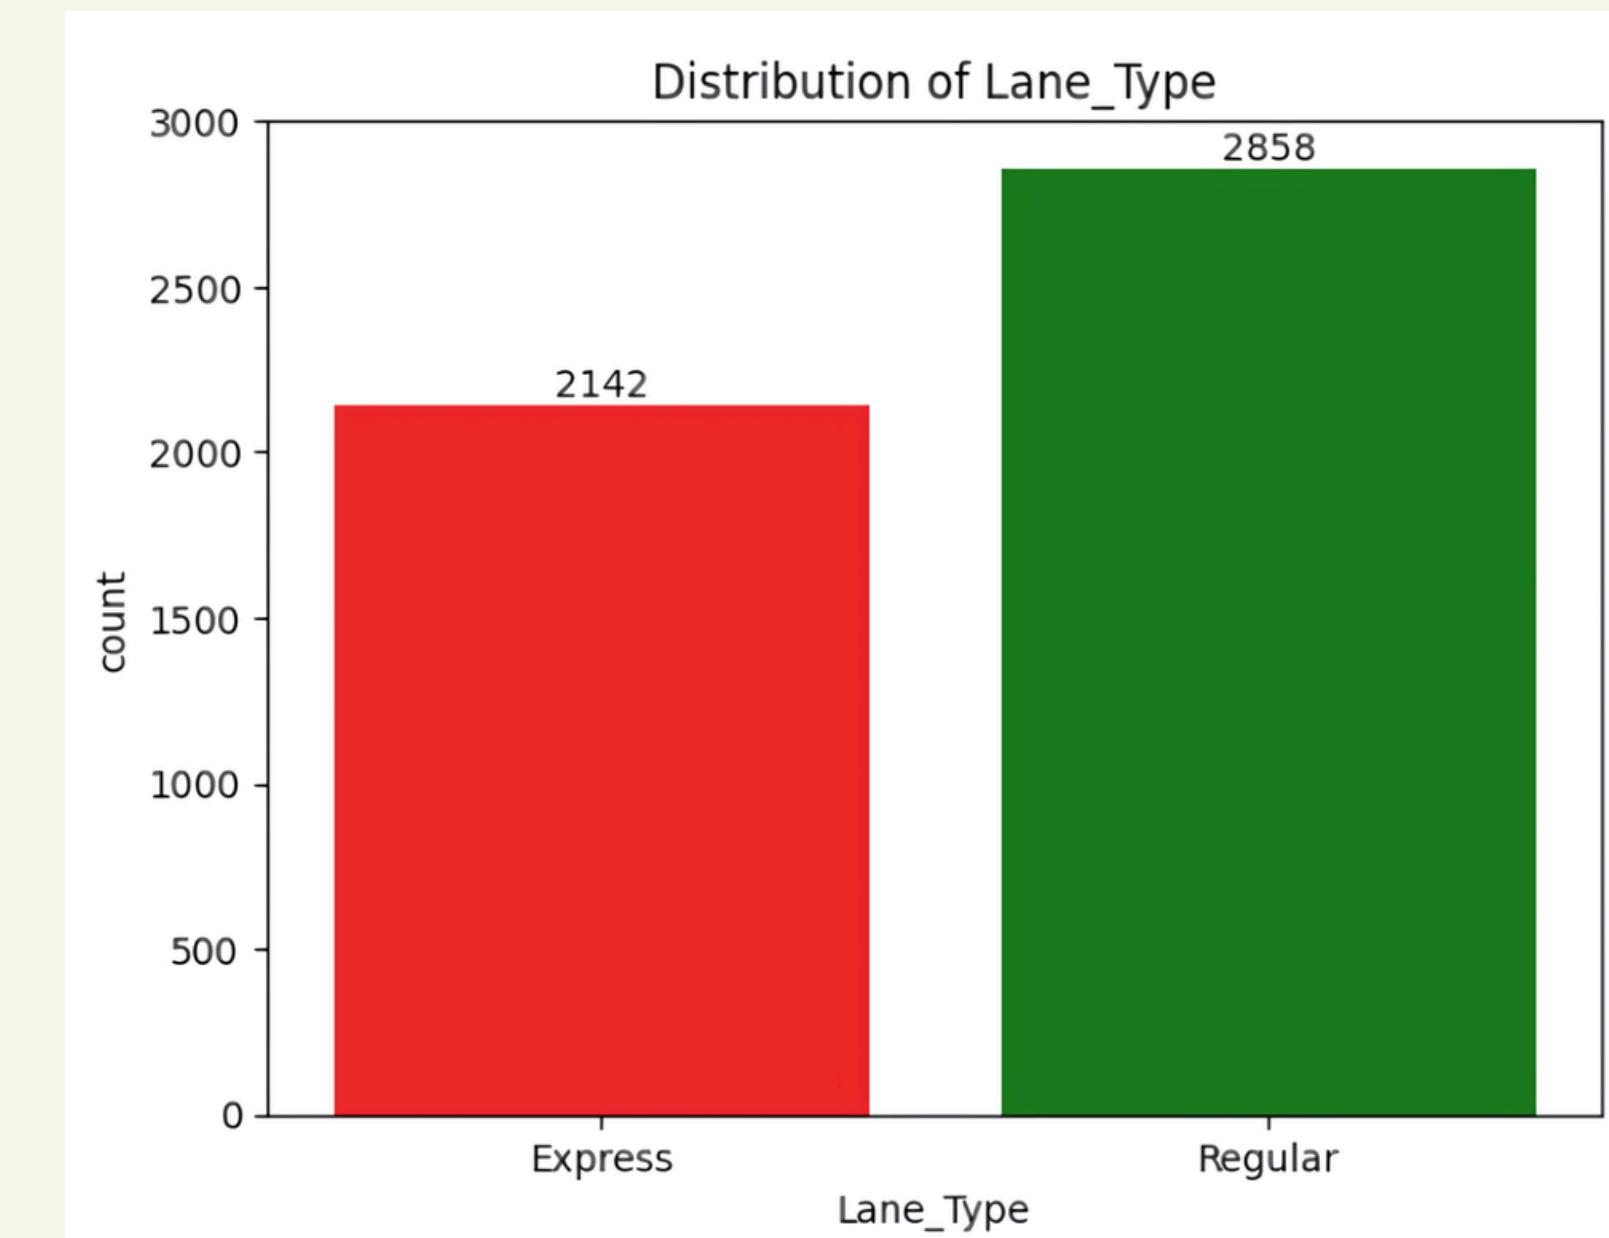


DATA VISUALIZATION

DISTRIBUTION OF VEHICLE TYPE

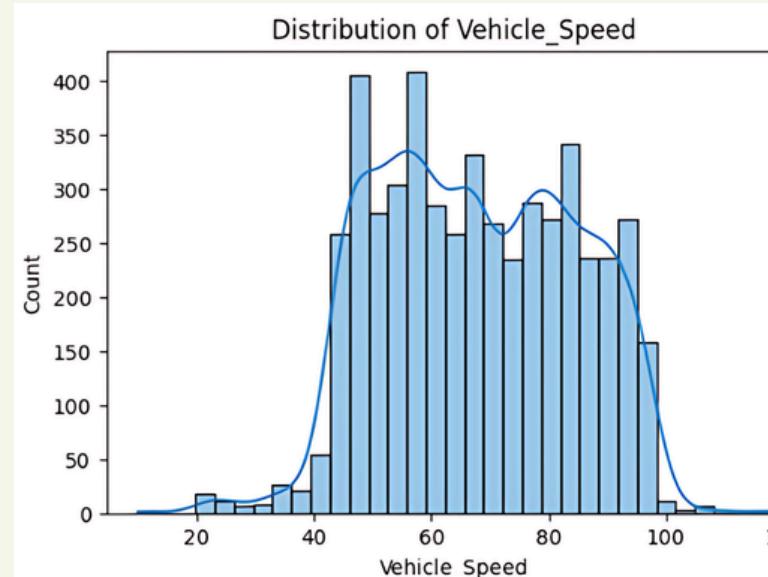
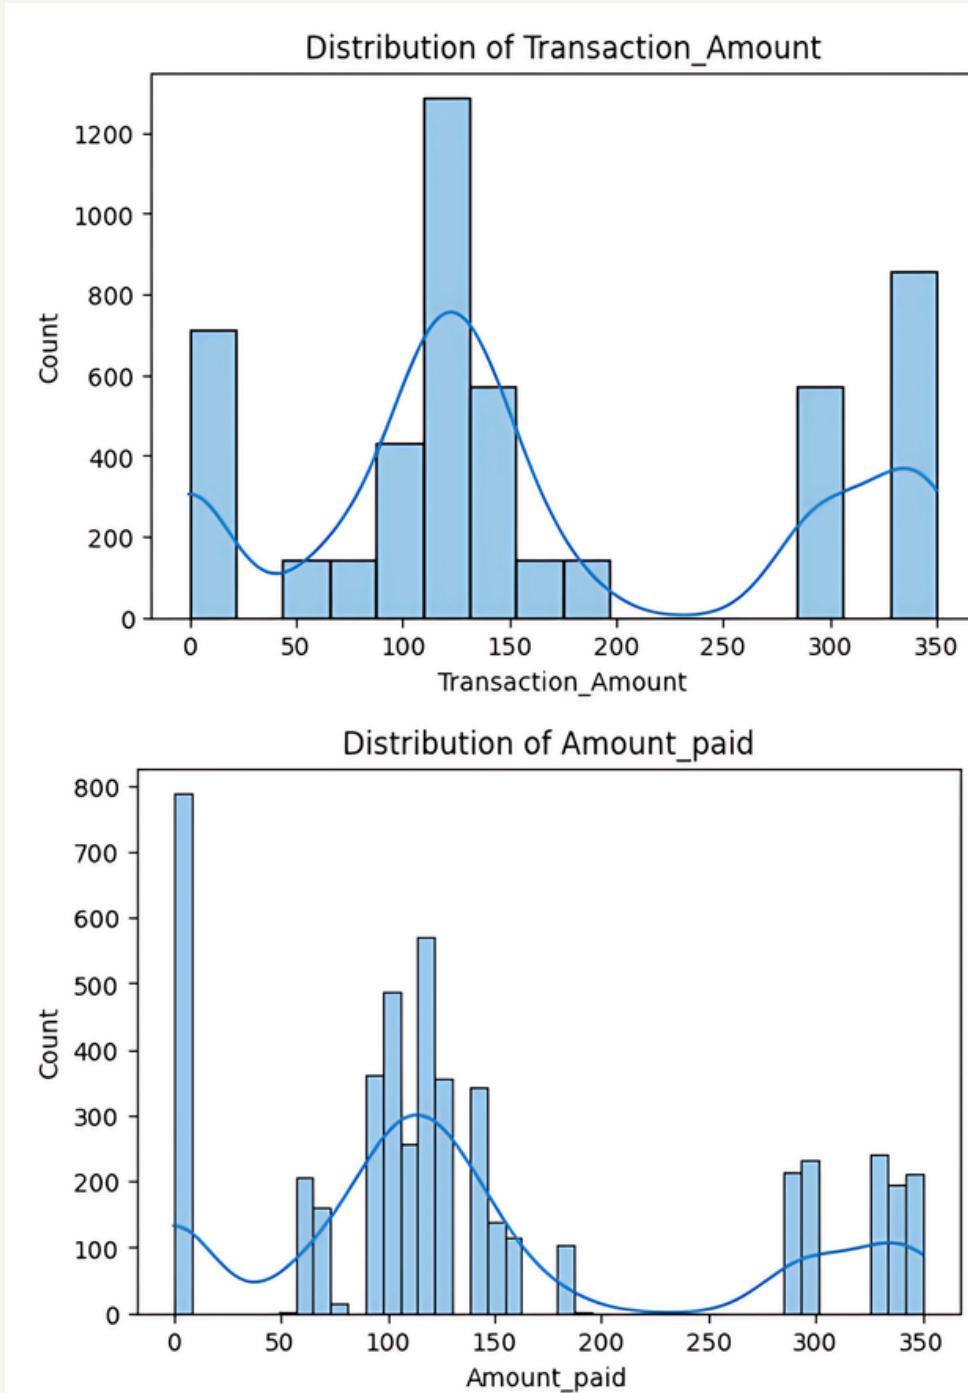


DISTRIBUTION OF LANE TYPE

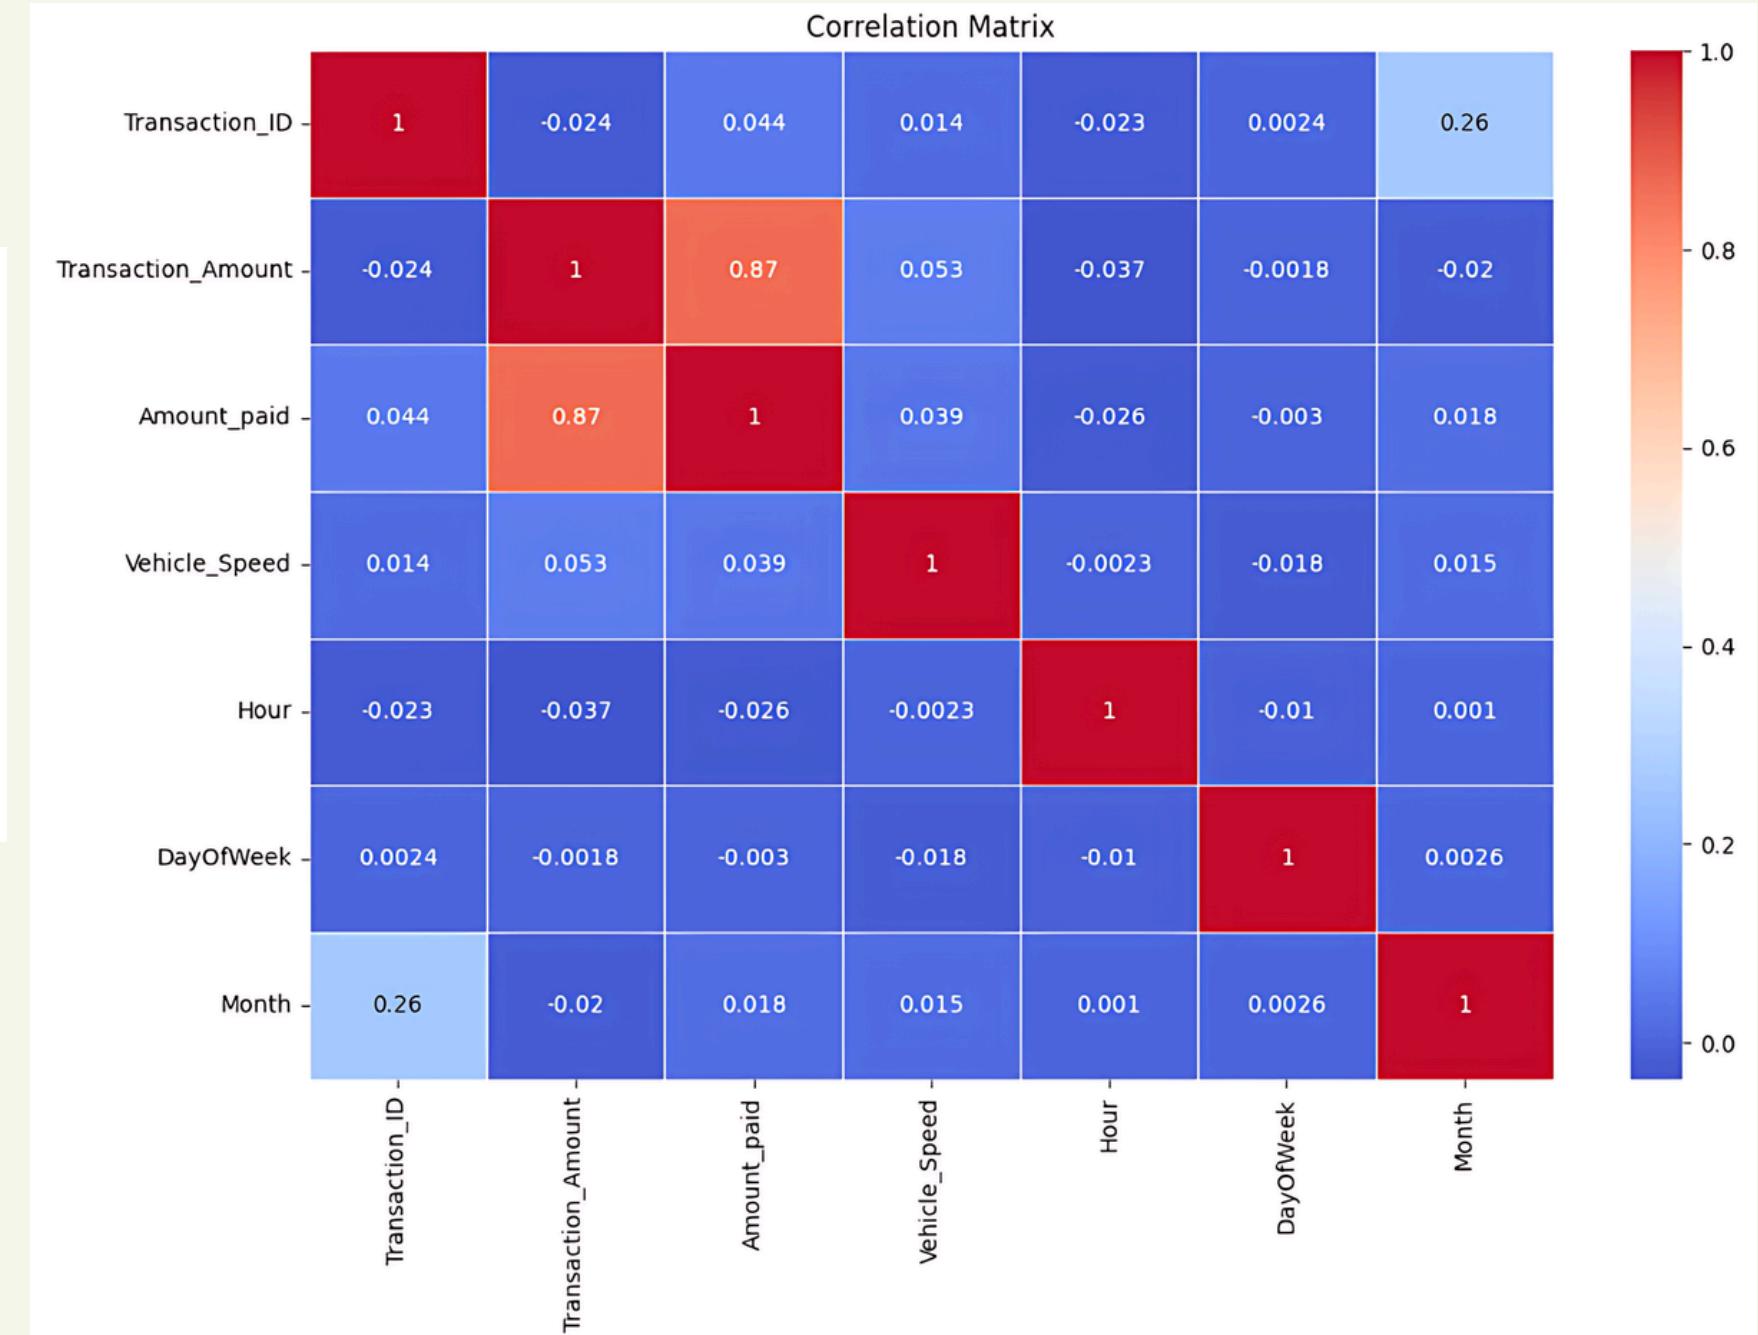


DATA VISUALIZATION

DISTRIBUTION OF NUMERICAL FEATURES



CORRELATION MATRIX



FEATURE ENGINEERING

```
# Encode categorical variables
categorical_features = ['Vehicle_Type', 'Lane_Type', 'Geographical_Location', 'Vehicle_Plate_Number','Vehicle_Dimensions', 'FastagID', 'TollBoothID','Fraud_indicator']
label_encoders = {}
for feature in categorical_features:
    label_encoders[feature] = LabelEncoder()
    df[feature] = label_encoders[feature].fit_transform(df[feature])

# Check the balance of the target variable
print("\nTarget variable distribution:")
print(df['Fraud_indicator'].value_counts())

# Address imbalance using SMOTE
X = df.drop(columns=['Transaction_ID', 'Timestamp', 'Fraud_indicator'])
y = df['Fraud_indicator']

smote = SMOTE(random_state=42)
X_res, y_res = smote.fit_resample(X, y)

print("\nBalanced target variable distribution:")
print(pd.Series(y_res).value_counts())
```

Target variable distribution:
1 4017
0 983
Name: Fraud_indicator, dtype: int64

Balanced target variable distribution:
0 4017
1 4017
Name: Fraud_indicator, dtype: int64

← Encode categorical variables using Label Encoding

← Checking Target Variable Balance

← Use SMOTE to handle class imbalance in the target variable.



Balancing Using SMOTE: Applies SMOTE to balance the classes by oversampling the minority class (fraudulent transactions) and undersampling the majority class (non-fraudulent transactions).

MODEL DEVELOPMENT & EVALUATION ...

SPLIT THE DATA INTO TRAINING AND TESTING SETS 70:30

```
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_res, y_res, test_size=0.3, random_state=42)

# Standardize the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

TRAIN MULTIPLE MODELS:

1. Random Forest
2. MLPClassifier (Neural Network)
3. Gradient Boosting
4. SVC (Support Vector Classifier)

FINE-TUNING:

Hyperparameter tuning using RandomizedSearchCV and GridSearchCV.

EVALUATION METRICS:

1. Accuracy
2. Precision
3. Recall
4. F1 Score
5. ROC-AUC Curve



MODEL DEVELOPMENT & EVALUATION

RANDOM FOREST CLASSIFIER

```
# Train a RandomForest classifier
clf = RandomForestClassifier(n_estimators=100, random_state=42)
clf.fit(X_train, y_train)

# Make predictions
y_pred = clf.predict(X_test)

# Evaluate the model
print("Random Forest Classifier")
print("\nClassification Report:")
report = classification_report(y_test, y_pred, digits=6)
print(report)

print("\nConfusion Matrix:")
conf_matrix = confusion_matrix(y_test, y_pred)
print(conf_matrix)

# Exact Calculation of Metrics
TP = conf_matrix[1, 1]
TN = conf_matrix[0, 0]
FP = conf_matrix[0, 1]
FN = conf_matrix[1, 0]

accuracy = (TP + TN) / (TP + TN + FP + FN)
precision = TP / (TP + FP)
recall = TP / (TP + FN)
f1_score = 2 * (precision * recall) / (precision + recall)

print(f"\nExact Accuracy: {accuracy:.10f}")
print(f"Exact Precision: {precision:.10f}")
print(f"Exact Recall: {recall:.10f}")
print(f"Exact F1 Score: {f1_score:.10f}")
```

Random Forest Classifier

Classification Report:

	precision	recall	f1-score	support
0	1.000000	0.986622	0.993266	1196
1	0.987002	1.000000	0.993459	1215
accuracy				0.993364
macro avg	0.993501	0.993311	0.993362	2411
weighted avg	0.993450	0.993364	0.993363	2411

Confusion Matrix:

```
[[1180  16]
 [   0 1215]]
```

Exact Accuracy: 0.9933637495

Exact Precision: 0.9870024370

Exact Recall: 1.0000000000

Exact F1 Score: 0.9934587081

0.9933

MODEL DEVELOPMENT & EVALUATION

RANDOM FOREST CLASSIFIER-HYPERTUNING

```
from sklearn.model_selection import RandomizedSearchCV
# Define the parameter grid
param_distributions = {
    'n_estimators': [100, 200, 300, 400, 500],
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth': [None, 10, 20, 30, 40, 50],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}

# Initialize RandomizedSearchCV
random_search = RandomizedSearchCV(estimator=clf, param_distributions=param_distributions,
                                    n_iter=100, cv=3, n_jobs=-1, verbose=2, random_state=42)

# Fit the random search to the data
random_search.fit(X_train, y_train)

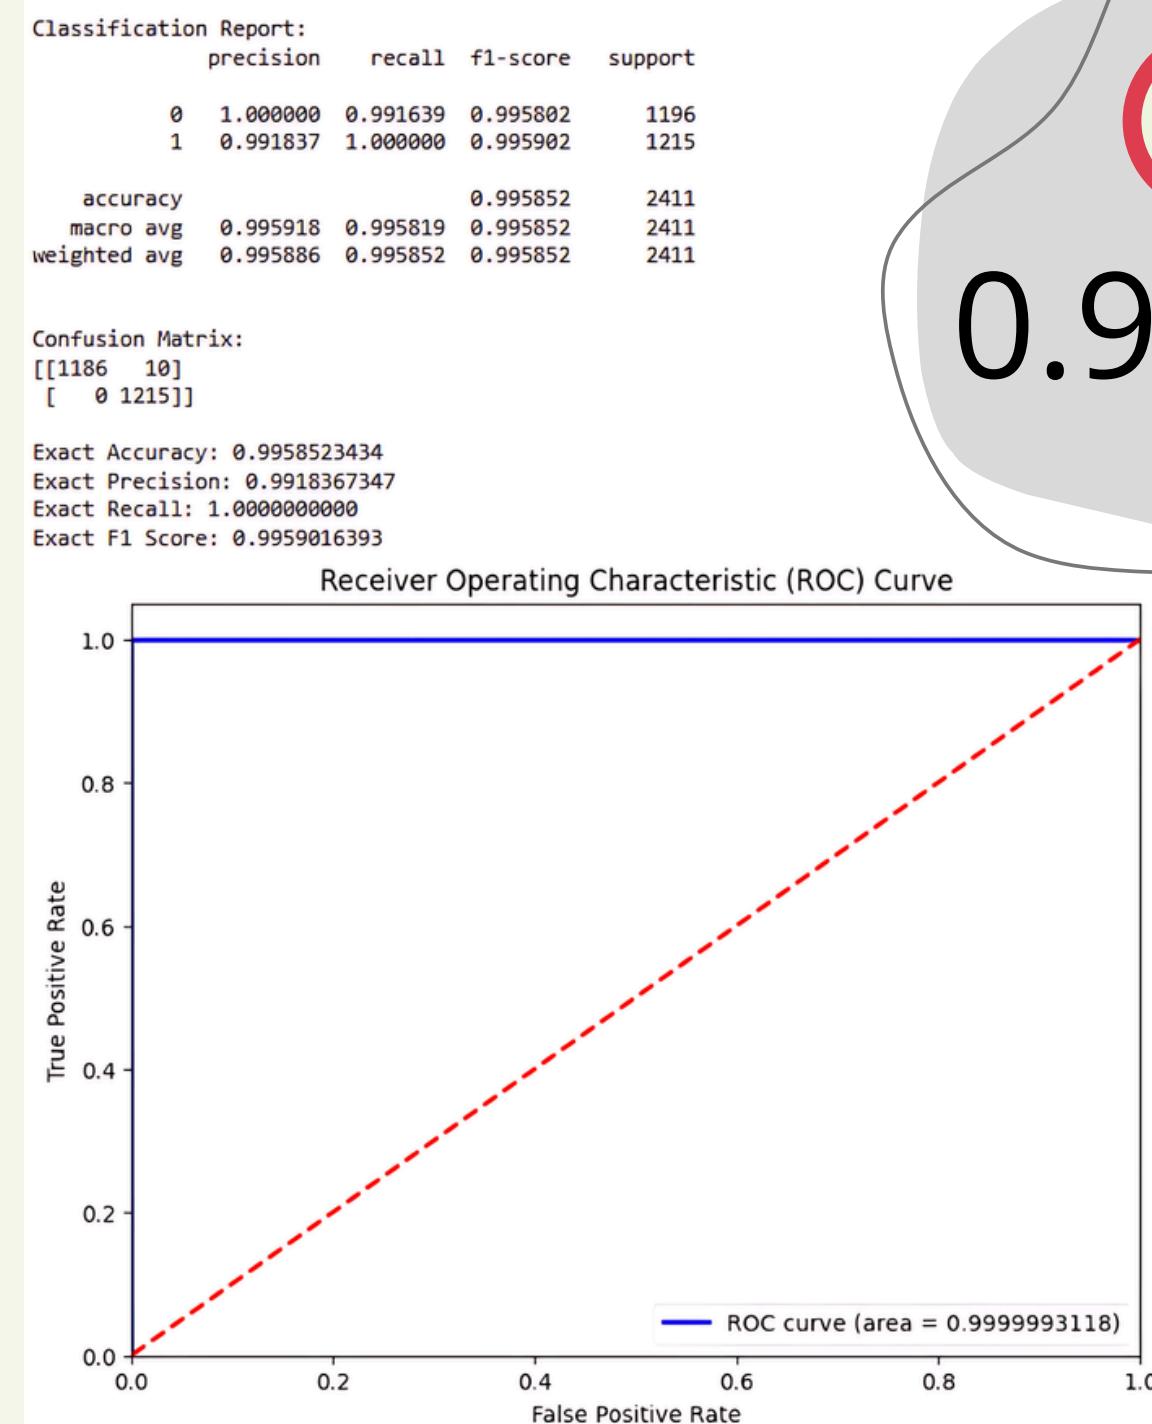
# Print the best parameters found by RandomizedSearchCV
print("Best parameters found: ", random_search.best_params_)

# Train a RandomForest classifier with the best parameters
best_clf = random_search.best_estimator_
best_clf.fit(X_train, y_train)

# Make predictions
y_pred = best_clf.predict(X_test)

# Evaluate the model
print("\nClassification Report:")
report = classification_report(y_test, y_pred, digits=6)
print(report)

print("\nConfusion Matrix:")
conf_matrix = confusion_matrix(y_test, y_pred)
print(conf_matrix)
```



MODEL DEVELOPMENT & EVALUATION

MULTI-LAYER PERCEPTRON (MLP)

```
# Initialize and train MLPClassifier (Neural Network)
from sklearn.neural_network import MLPClassifier
mlp_clf = MLPClassifier(hidden_layer_sizes=(100, 50), max_iter=100, random_state=42)
mlp_clf.fit(X_train, y_train)

# Make predictions
mlp_y_pred = mlp_clf.predict(X_test)

# Evaluate MLPClassifier
print("\nMLPClassifier - Classification Report:")
print(classification_report(y_test, mlp_y_pred, digits=6))

print("\nMLPClassifier - Confusion Matrix:")
conf_matrix1=confusion_matrix(y_test, mlp_y_pred)
print(conf_matrix1)

# Exact Calculation of Metrics
TP = conf_matrix1[1, 1]
TN = conf_matrix1[0, 0]
FP = conf_matrix1[0, 1]
FN = conf_matrix1[1, 0]

accuracy = (TP + TN) / (TP + TN + FP + FN)
precision = TP / (TP + FP)
recall = TP / (TP + FN)
f1_score = 2 * (precision * recall) / (precision + recall)

print(f"\nExact Accuracy: {accuracy:.10f}")
print(f"Exact Precision: {precision:.10f}")
print(f"Exact Recall: {recall:.10f}")
print(f"Exact F1 Score: {f1_score:.10f}")
```

MLPClassifier - Classification Report:				
	precision	recall	f1-score	support
0	0.980000	0.983278	0.981636	1196
1	0.983485	0.980247	0.981863	1215
accuracy			0.981750	2411
macro avg	0.981742	0.981762	0.981750	2411
weighted avg	0.981756	0.981750	0.981750	2411

MLPClassifier - Confusion Matrix:
[[1176 20]
 [24 1191]]

Exact Accuracy: 0.9817503111
Exact Precision: 0.9834847234
Exact Recall: 0.9802469136
Exact F1 Score: 0.9818631492

0.9817

MODEL DEVELOPMENT & EVALUATION

Multi-Layer Perceptron (MLP)-HYPERTUNING

```
# Define the parameter grid
param_distributions = {
    'hidden_layer_sizes': [(50,), (100,), (100, 50), (100, 100, 50)],
    'activation': ['tanh', 'relu'],
    'solver': ['sgd', 'adam'],
    'alpha': [0.0001, 0.001, 0.01, 0.1],
    'learning_rate': ['constant', 'adaptive'],
    'max_iter': [200, 300, 400, 500]
}

# Initialize the MLPClassifier
mlp_clf = MLPClassifier(random_state=42)

# Initialize RandomizedSearchCV
random_search = RandomizedSearchCV(estimator=mlp_clf, param_distributions=param_distributions,
                                    n_iter=50, cv=3, n_jobs=-1, verbose=2, random_state=42)

# Fit the random search to the data
random_search.fit(X_train, y_train)

# Print the best parameters found by RandomizedSearchCV
print("Best parameters found: ", random_search.best_params_)

# Train a MLPClassifier with the best parameters
best_mlp_clf = random_search.best_estimator_
best_mlp_clf.fit(X_train, y_train)

# Make predictions
mlp_y_pred = best_mlp_clf.predict(X_test)

# Evaluate the model
print("\nMLPClassifier - Classification Report:")
report = classification_report(y_test, mlp_y_pred, digits=6)
print(report)

print("\nMLPClassifier - Confusion Matrix:")
conf_matrix = confusion_matrix(y_test, mlp_y_pred)
print(conf_matrix)
```

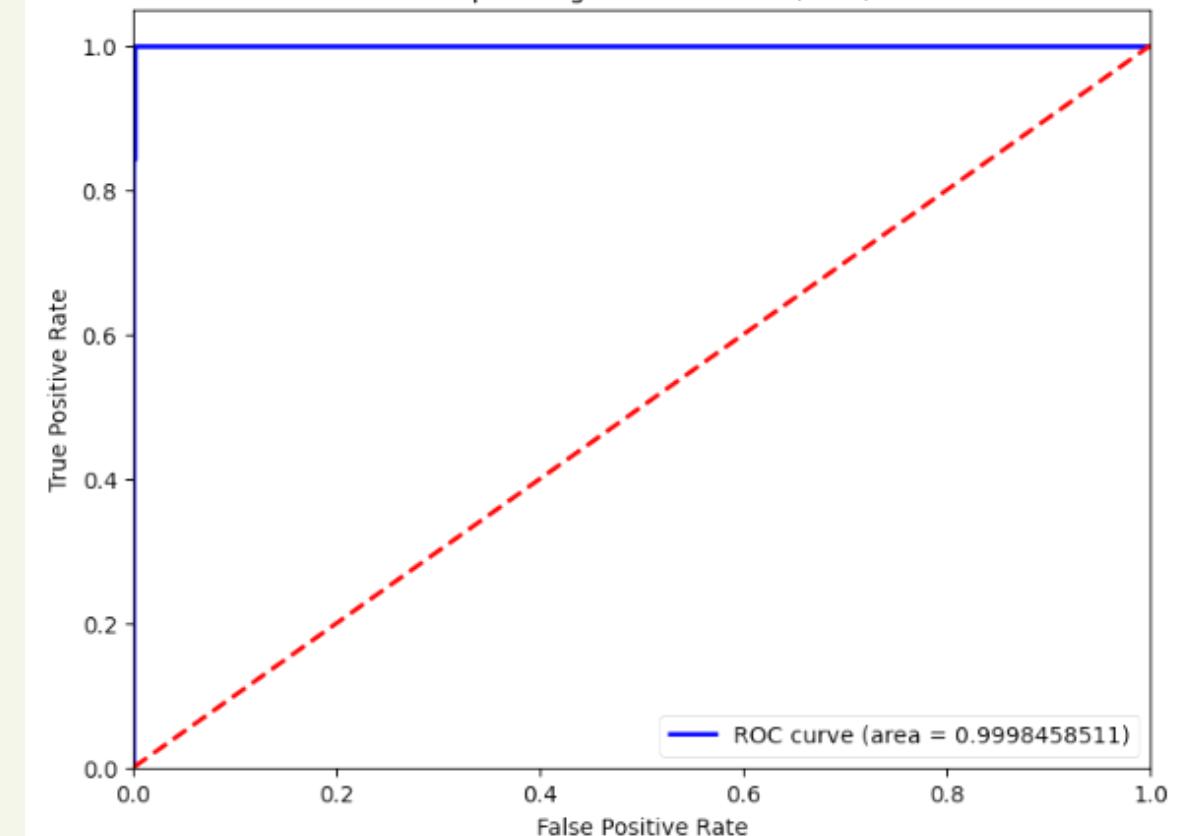
```
MLPClassifier - Classification Report:
precision    recall   f1-score  support
          0   1.000000  0.994983  0.997485   1196
          1   0.995086  1.000000  0.997537   1215

   accuracy                           0.997511   2411
    macro avg   0.997543  0.997492  0.997511   2411
weighted avg   0.997524  0.997511  0.997511   2411
```

```
MLPClassifier - Confusion Matrix:
[[1190   6]
 [  0 1215]]

Exact Accuracy: 0.9975114061
Exact Precision: 0.9950859951
Exact Recall: 1.0000000000
Exact F1 Score: 0.9975369458
```

Receiver Operating Characteristic (ROC) Curve



MODEL DEVELOPMENT & EVALUATION

GRADIENT BOOSTING CLASSIFIER

```
# Initialize and train GradientBoostingClassifier
gb_clf = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, random_state=42)
gb_clf.fit(X_train, y_train)

# Make predictions
gb_y_pred = gb_clf.predict(X_test)

# Evaluate GradientBoostingClassifier
print("\nGradientBoostingClassifier - Classification Report:")
print(classification_report(y_test, gb_y_pred, digits=6))

print("\nGradientBoostingClassifier - Confusion Matrix:")
conf_matrix2=confusion_matrix(y_test, gb_y_pred)
print(conf_matrix2)

# Exact Calculation of Metrics
TP = conf_matrix2[1, 1]
TN = conf_matrix2[0, 0]
FP = conf_matrix2[0, 1]
FN = conf_matrix2[1, 0]

accuracy = (TP + TN) / (TP + TN + FP + FN)
precision = TP / (TP + FP)
recall = TP / (TP + FN)
f1_score = 2 * (precision * recall) / (precision + recall)

print(f"\nExact Accuracy: {accuracy:.10f}")
print(f"Exact Precision: {precision:.10f}")
print(f"Exact Recall: {recall:.10f}")
print(f"Exact F1 Score: {f1_score:.10f}")
```

GradientBoostingClassifier - Classification Report:

	precision	recall	f1-score	support
0	1.000000	0.988294	0.994113	1196
1	0.988609	1.000000	0.994272	1215
accuracy			0.994193	2411
macro avg	0.994304	0.994147	0.994192	2411
weighted avg	0.994259	0.994193	0.994193	2411

GradientBoostingClassifier - Confusion Matrix:
[[1182 14]
 [0 1215]]

Exact Accuracy: 0.9941932808
Exact Precision: 0.9886086249
Exact Recall: 1.0000000000
Exact F1 Score: 0.9942716858

0.9941

MODEL DEVELOPMENT & EVALUATION

GRADIENT BOOSTING CLASSIFIER- HYPERPARAMETER TUNING

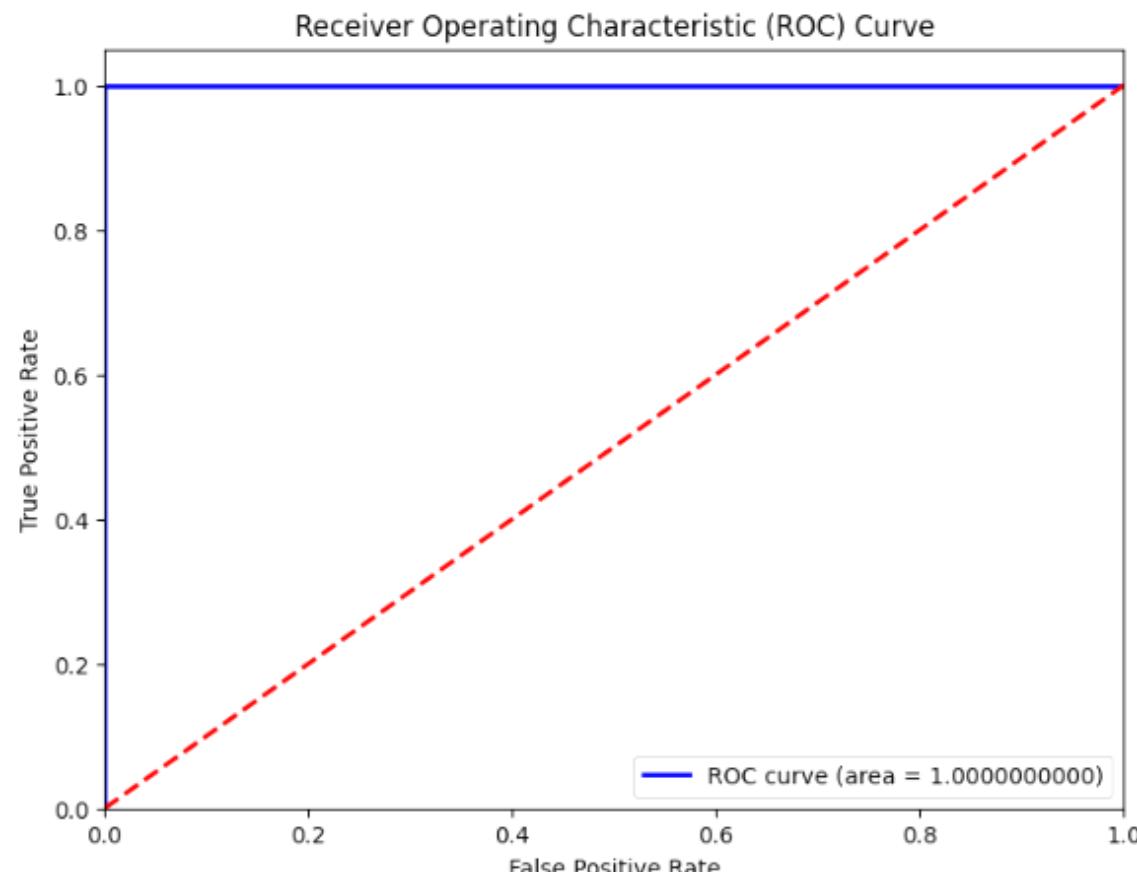
```
param_distributions = {  
    'n_estimators': [100, 200, 300, 400, 500],  
    'learning_rate': [0.01, 0.05, 0.1, 0.2, 0.3],  
    'max_depth': [3, 5, 7, 9],  
    'min_samples_split': [2, 5, 10],  
    'min_samples_leaf': [1, 2, 4],  
    'subsample': [0.8, 0.9, 1.0],  
    'max_features': ['auto', 'sqrt', 'log2', None]  
}  
  
# Initialize the GradientBoostingClassifier  
gb_clf = GradientBoostingClassifier(random_state=42)  
  
# Initialize RandomizedSearchCV  
random_search = RandomizedSearchCV(estimator=gb_clf, param_distributions=param_distributions,  
                                    n_iter=100, cv=3, n_jobs=-1, verbose=2, random_state=42)  
  
# Fit the random search to the data  
random_search.fit(X_train, y_train)  
  
# Print the best parameters found by RandomizedSearchCV  
print("Best parameters found: ", random_search.best_params_)  
  
# Train a GradientBoostingClassifier with the best parameters  
best_gb_clf = random_search.best_estimator_  
best_gb_clf.fit(X_train, y_train)  
  
# Make predictions  
gb_y_pred = best_gb_clf.predict(X_test)  
  
# Evaluate the model  
print("\nGradientBoostingClassifier - Classification Report:")  
report = classification_report(y_test, gb_y_pred, digits=6)  
print(report)  
  
print("\nGradientBoostingClassifier - Confusion Matrix:")  
conf_matrix = confusion_matrix(y_test, gb_y_pred)  
print(conf_matrix)  
  
# Exact Calculation of Metrics  
TP = conf_matrix[1, 1]  
TN = conf_matrix[0, 0]  
FP = conf_matrix[0, 1]  
FN = conf_matrix[1, 0]  
  
accuracy = (TP + TN) / (TP + TN + FP + FN)  
precision = TP / (TP + FP)  
recall = TP / (TP + FN)  
f1_score = 2 * (precision * recall) / (precision + recall)
```

GradientBoostingClassifier - Classification Report:
precision recall f1-score support
0 1.000000 0.998328 0.999163 1196
1 0.998357 1.000000 0.999178 1215

accuracy 0.999170 2411
macro avg 0.999178 0.999164 0.999170 2411
weighted avg 0.999172 0.999170 0.999170 2411

GradientBoostingClassifier - Confusion Matrix:
[[1194 2]
 [0 1215]]

Exact Accuracy: 0.9991704687
Exact Precision: 0.9983566146
Exact Recall: 1.0000000000
Exact F1 Score: 0.9991776316



0.9991

MODEL DEVELOPMENT & EVALUATION

SUPPORT VECTOR CLASSIFIER

```
# Initialize and train SVC (Support Vector Classifier)
svm_clf = SVC(kernel='rbf', probability=True, random_state=42)
svm_clf.fit(X_train, y_train)

# Make predictions
svm_y_pred = svm_clf.predict(X_test)

# Evaluate SVC
print("\nSVC - Classification Report:")
print(classification_report(y_test, svm_y_pred, digits=6))

print("\nSVC - Confusion Matrix:")
conf_matrix3 = confusion_matrix(y_test, svm_y_pred)
print(conf_matrix3)

# Exact Calculation of Metrics
TP = conf_matrix3[1, 1]
TN = conf_matrix3[0, 0]
FP = conf_matrix3[0, 1]
FN = conf_matrix3[1, 0]

accuracy = (TP + TN) / (TP + TN + FP + FN)
precision = TP / (TP + FP)
recall = TP / (TP + FN)
f1_score = 2 * (precision * recall) / (precision + recall)

print(f"\nExact Accuracy: {accuracy:.10f}")
print(f"Exact Precision: {precision:.10f}")
print(f"Exact Recall: {recall:.10f}")
print(f"Exact F1 Score: {f1_score:.10f}")
```

SVC - Classification Report:				
	precision	recall	f1-score	support
0	0.994686	0.938963	0.966022	1196
1	0.943058	0.995062	0.968362	1215
accuracy			0.967234	2411
macro avg	0.968872	0.967012	0.967192	2411
weighted avg	0.968668	0.967234	0.967201	2411

SVC - Confusion Matrix:
[[1123 73]
 [6 1209]]

Exact Accuracy: 0.9672335131
Exact Precision: 0.9430577223
Exact Recall: 0.9950617284
Exact F1 Score: 0.9683620344

0.9672

MODEL DEVELOPMENT & EVALUATION

SUPPORT VECTOR CLASSIFIER- HYPERTUNING

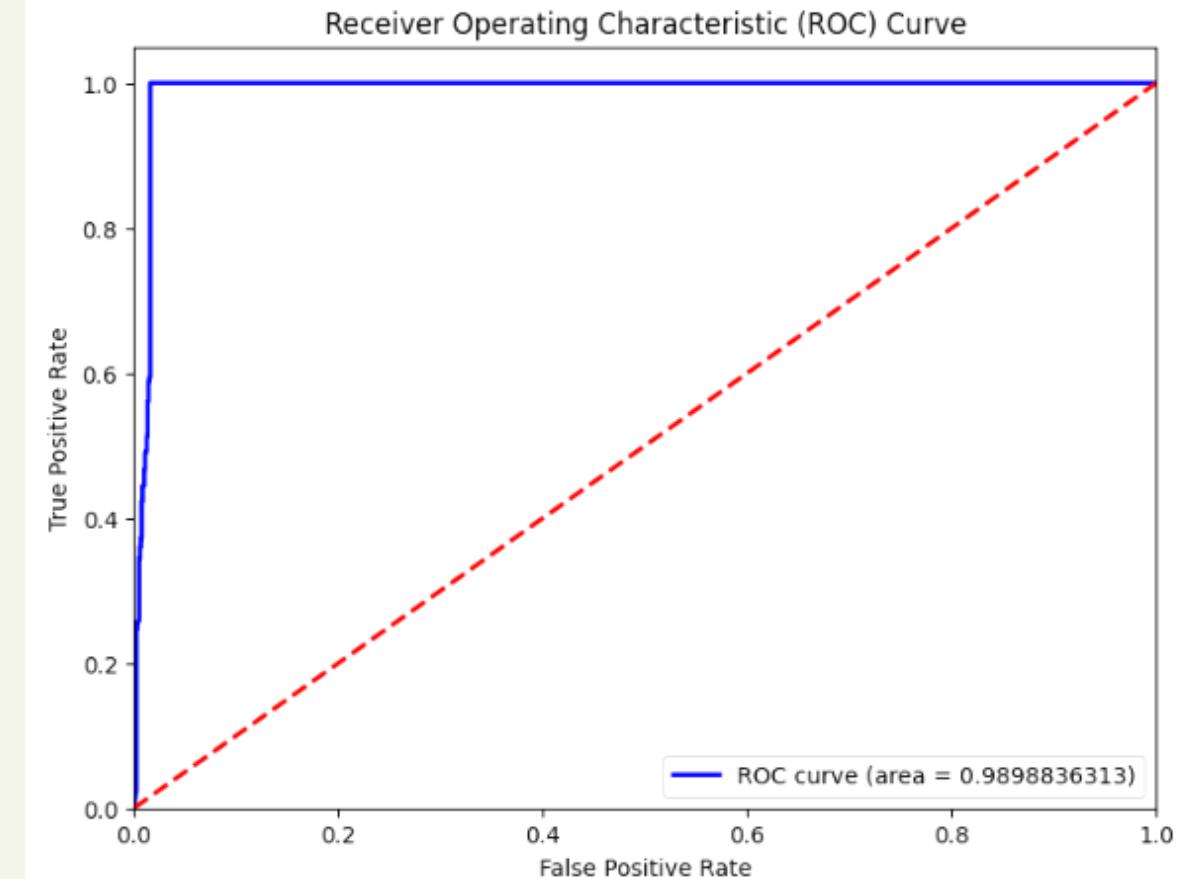
```
param_grid = {  
    'C': [0.1, 1, 10, 100],  
    'gamma': [1, 0.1, 0.01, 0.001],  
    'kernel': ['rbf']  
}  
  
# Initialize the SVC  
svm_clf = SVC(probability=True, random_state=42)  
  
# Initialize GridSearchCV  
grid_search = GridSearchCV(estimator=svm_clf, param_grid=param_grid,  
                           cv=3, n_jobs=-1, verbose=2)  
  
# Fit GridSearchCV to the data  
grid_search.fit(X_train, y_train)  
  
# Print the best parameters found by GridSearchCV  
print("Best parameters found: ", grid_search.best_params_)  
  
# Train a SVC with the best parameters  
best_svm_clf = grid_search.best_estimator_  
best_svm_clf.fit(X_train, y_train)  
  
# Make predictions  
svm_y_pred = best_svm_clf.predict(X_test)  
  
# Evaluate the model  
print("\nSVC - Classification Report:")  
report = classification_report(y_test, svm_y_pred, digits=6)  
print(report)  
  
print("\nSVC - Confusion Matrix:")  
conf_matrix = confusion_matrix(y_test, svm_y_pred)  
print(conf_matrix)  
  
# Exact Calculation of Metrics  
TP = conf_matrix[1, 1]  
TN = conf_matrix[0, 0]  
FP = conf_matrix[0, 1]  
FN = conf_matrix[1, 0]  
  
accuracy = (TP + TN) / (TP + TN + FP + FN)  
precision = TP / (TP + FP)  
recall = TP / (TP + FN)  
f1_score = 2 * (precision * recall) / (precision + recall)
```

SVC - Classification Report:

	precision	recall	f1-score	support
0	1.000000	0.976589	0.988156	1196
1	0.977474	1.000000	0.988609	1215
accuracy			0.988387	2411
macro avg	0.988737	0.988294	0.988382	2411
weighted avg	0.988648	0.988387	0.988384	2411

SVC - Confusion Matrix:
[[1168 28]
 [0 1215]]

Exact Accuracy: 0.9883865616
Exact Precision: 0.9774738536
Exact Recall: 1.0000000000
Exact F1 Score: 0.9886086249



0.9883

MODEL SELECTION

SELECTED MODEL: GRADIENT BOOSTING CLASSIFIER

REASON:

This model was selected based on its superior performance across the evaluation metrics. The Gradient Boosting model, fine-tuned with Randomized Search, achieved near-perfect performance with accuracy(0.9991). Ensures high precision and recall, critical for minimizing false negatives and maintaining transaction integrity.

- **Exact Metrics:**

- Accuracy: 0.9991704687
- Precision: 0.9983566146
- Recall: 1.0000000000
- F1 Score: 0.9991776316

- **ROC Curve:**

- AUC: 1.0000
- Demonstrates excellent discrimination between fraudulent and non-fraudulent transactions



REAL TIME FRAUD DETECTION

DEPLOYMENT ON FLASK

The screenshot shows a Microsoft Visual Studio Code (VS Code) interface. The title bar reads "Project-2.FastagFraudDetection". The left sidebar contains icons for Explorer, Search, Open, and Python. The Explorer view shows a project structure under "PROJECT-2.FA...":

- FastagFraudDetection
 - static
 - fastag.jpg
 - left-img.jpg
 - right-img.jpg
 - templates
 - index.html
 - app.py
 - gradient_boosting_model.pkl

The main editor tab is "app.py", which contains the following code:

```
1 import numpy as np
2 from flask import Flask, request, jsonify, render_template
3 from sklearn.preprocessing import LabelEncoder
4 import joblib
5
6 app = Flask(__name__)
7
8 # Load the pre-trained model
9 model = joblib.load('gradient_boosting_model.pkl')
10
11 # Define categories for each categorical feature
12 categories = {
13     'Vehicle_Type': ['Bus', 'Car', 'Motorcycle', 'SUV', 'Sedan', 'Truck', 'Van'],
14     'Lane_Type': ['Regular', 'Express'],
15     'Geographical_Location': [
16         '13.059816123454882, 77.77068662374292',
17         '13.042660878688794, 77.47580097259879',
18         '12.84197701525119, 77.67547528176169',
19         '12.936687032945434, 77.53113977439017',
20         '13.21331620748757, 77.55413526894684'], # Add predefined categories or handle free
21     'Vehicle_Plate_Number': [], # Add predefined categories or handle free text as needed
22     'Vehicle_Dimensions': ['Large', 'Medium', 'Small'], # Add predefined categories or handle
23     'FastagID': [], # Add predefined categories or handle free text as needed
24     'TollBoothID': ['B-102', 'A-101', 'C-103', 'D-106', 'D-105', 'D-104'] # Example toll bo
25 }
```

The bottom navigation bar includes "PROBLEMS", "OUTPUT", "DEBUG CONSOLE", "TERMINAL" (which is selected), and "PORTS". The terminal window shows the following log output:

```
127.0.0.1 - - [22/Jun/2024 01:45:12] "GET /static/left-img.jpg HTTP/1.1" 304 -
127.0.0.1 - - [22/Jun/2024 01:45:12] "GET /static/right.jpg HTTP/1.1" 404 -
* History restored
```

The status bar at the bottom indicates the path: "PS C:\Users\Sonal R Sonarghare\Documents\Documents\Internship\Mentorness\Project-2.FastagFraudDetection".

REAL TIME FRAUD DETECTION



Fraud Detection

Vehicle Type:

Lane Type:

Geographical Location:

Vehicle Plate Number:

Vehicle Dimensions:

Fastag ID:

TollBooth ID:

Transaction Amount:

Amount Paid:

Vehicle Speed:

Hour:

DayOfWeek:

Month:

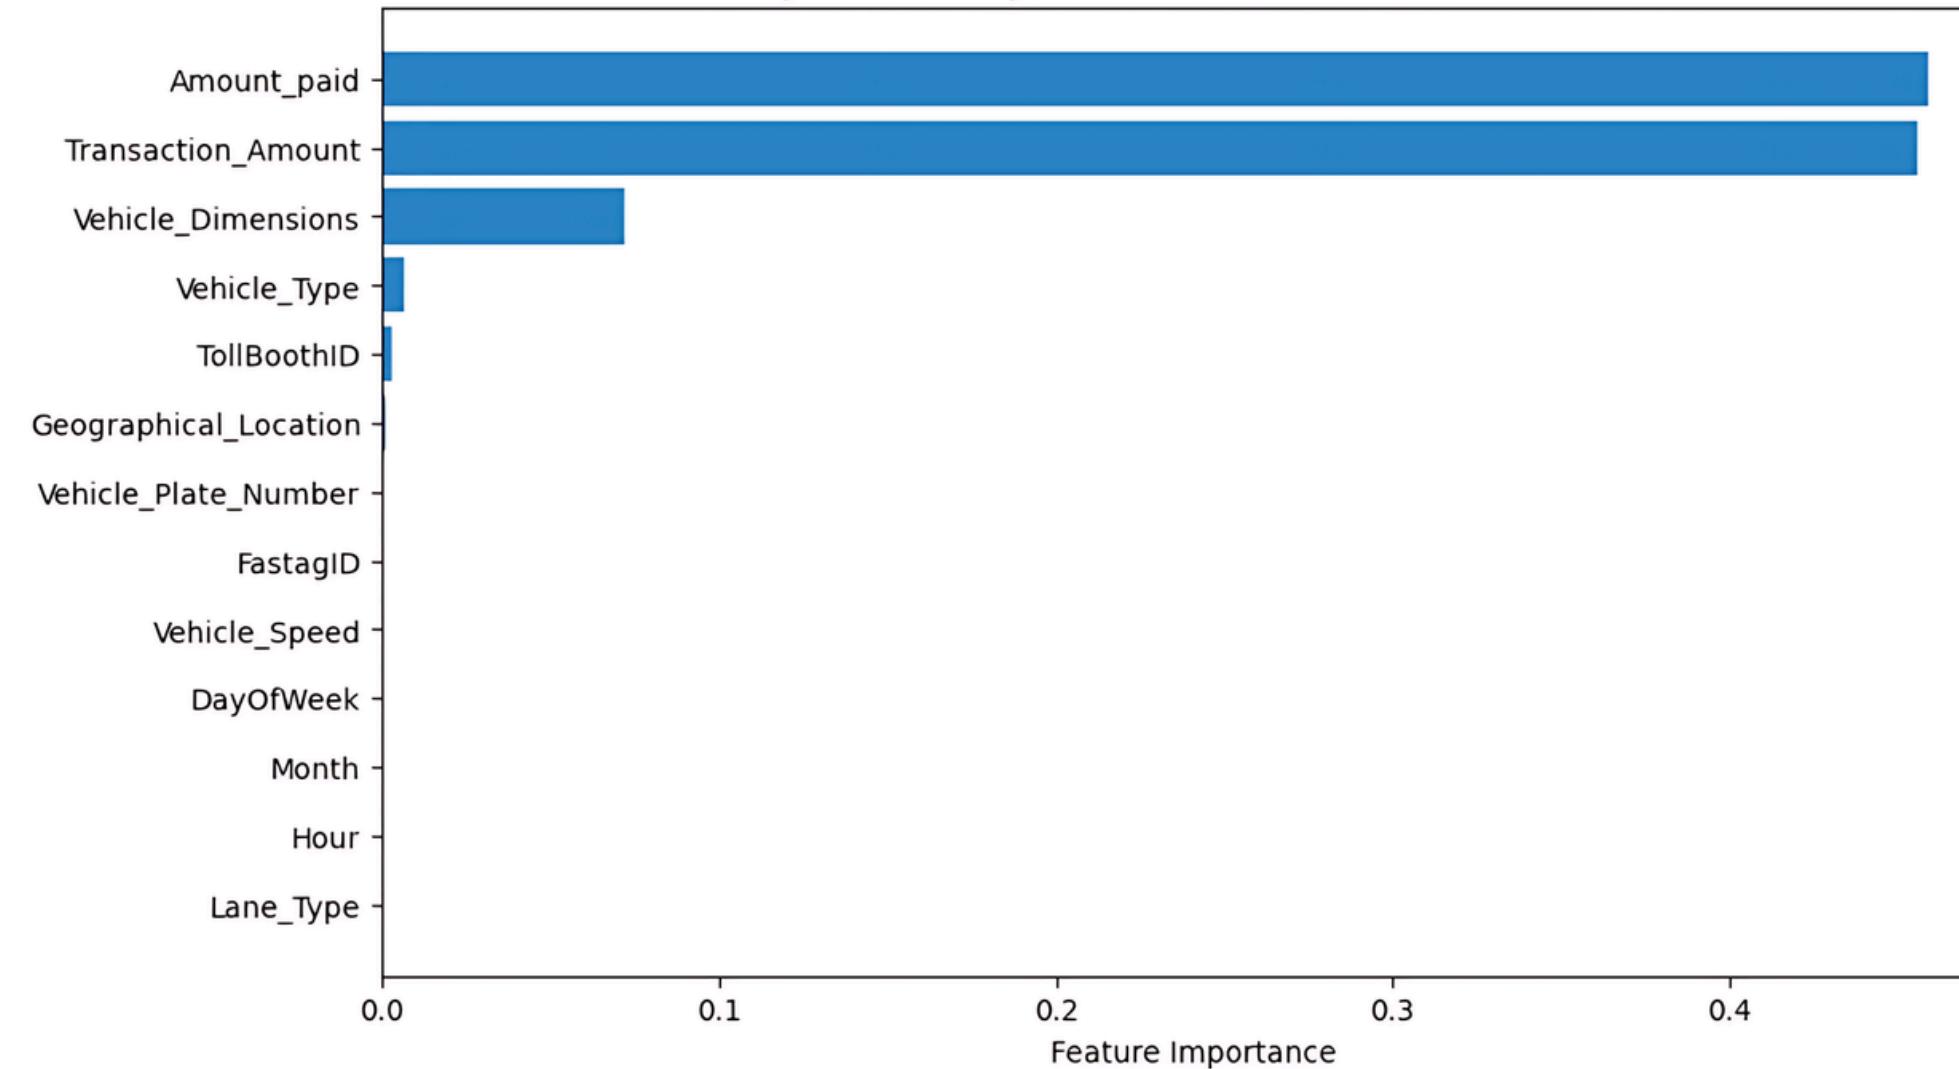
Fraud Indicator: Fraud
Transaction is fraudulent



EXPLANTORY ANALYSIS

Amount_paid: 0.4593686619
Transaction_Amount: 0.4562961305
Vehicle_Dimensions: 0.0719769523
Vehicle_Type: 0.0069367856
TollBoothID: 0.0030344974
Geographical_Location: 0.0012082614
Vehicle_Plate_Number: 0.0005814439
FastagID: 0.0004597983
Vehicle_Speed: 0.0000912652
DayOfWeek: 0.0000175068
Month: 0.0000175064
Hour: 0.0000111883
Lane_Type: 0.0000000021

Top 13 Most Important Features for Fraud Detection



The chart shows that the most critical features for fraud detection are:

Amount_paid (0.459)

Transaction_Amount (0.456)

Vehicle_Dimensions (0.072)

These three features have the highest importance scores, indicating they are the most influential in identifying fraud. Other features, such as Vehicle_Type, TollBoothID, and Geographical_Location, have minimal importance, while features like Vehicle_Speed, DayOfWeek, Month, Hour, and Lane_Type contribute very little to fraud detection.

CONCLUSION

Model Selected:

Gradient Boosting Classifier

Performance:

Achieved near-perfect accuracy (99.92%) in detecting fraudulent FASTag transactions.

High precision (99.84%) and recall (100%) rates ensure minimal false positives and negatives.

Why Gradient Boosting:

- Superior in handling complex patterns and imbalanced data.
- Effective feature importance analysis.
- Proven success in similar financial fraud detection applications.

Impact:

- Enhanced security and integrity of FASTag transactions.
- Significant reduction in potential revenue losses due to fraud.

Future Work:

- Continuous model monitoring and updates with new data.
- Explore additional features and advanced ensemble techniques for further improvements.



#GoCashLess



THANK YOU

-  SONAL SONARGHARE
-  sonalsonarghare30@gmail.com
-  A.P SHAH INSTITUTE OF TECHNOLOGY

